Minimum Wages and Employment: A Replication and Extension of Card & Krueger (1994)

Course: ECON6067 · Quantitative Tools for Economic Analysis

Project: 01 — Card & Krueger (1994), Minimum Wage / Difference-in-Differences

Replication target: Tables 2–4 and Figure 1 (Table 1 omitted: its sample-frame data are restricted)

---

1. Research question and motivation

On 1 April 1992, New Jersey raised its state minimum wage from \$4.25 to \$5.05 per hour — an increase of roughly 19 percent — while the neighbouring state of Pennsylvania, separated from New Jersey only by the Delaware River, retained the federal minimum wage of \$4.25. The prediction of the standard competitive model of the labour market is unambiguous and dates back at least to Stigler (1946): a binding minimum wage raises the price of labour above its market-clearing level, so profit-maximising firms should employ fewer low-wage workers. The magnitude of this response, however, has been the subject of one of the longest-running and most consequential debates in empirical labour economics.

The early time-series literature, surveyed by Brown, Gilroy, and Kohen (1982, 1983), concluded that a 10-percent increase in the minimum wage reduces teenage employment by roughly 1 to 3 percent — small but negative effects, consistent with the competitive prediction. Wellington (1991) reached a similar conclusion. Yet a wave of establishment-level studies that emerged in the early 1990s, using carefully constructed comparison groups rather than aggregate time series, began to tell a different story. Katz and Krueger (1992) found that the 1990–1991 federal minimum-wage increases had no adverse effect on employment in Texas fast-food restaurants; Card (1992a, 1992b) reported similar results for California and the federal changes. These studies raised the possibility that the conventional wisdom rested on fragile identification rather than a genuine employment effect.

Card and Krueger (1994) — hereafter CK — provide the most celebrated and influential piece of this new evidence. They exploit New Jersey's April 1992 minimum-wage increase as a quasi-natural experiment. Their design surveys the same 410 fast-food restaurants in New Jersey and eastern Pennsylvania twice: once in February–March 1992, a little over a month before the increase, and again in November–December 1992, seven to eight months after. The central research question of the paper, and of this replication, is therefore:

> **Did the April 1992 New Jersey minimum-wage increase reduce employment in the fast-food industry?**

The paper's answer is striking: employment in New Jersey *increased* relative to Pennsylvania. The difference-in-differences (DID) estimate is about +2.75 full-time-equivalent (FTE) employees per store, with a t-statistic just above 2. Far from confirming the competitive prediction of employment losses, the paper finds no evidence of any decline.

Three features of the setting make it especially valuable for teaching and for replication. First, the increase occurred during a regional recession, so it is unlikely that the treatment effect is confounded with a coincident boom in New Jersey. Second, New Jersey and eastern Pennsylvania are tightly linked labour and product markets, which makes Pennsylvania a natural control group; at the same time, within New Jersey there is substantial cross-store variation in the initial starting wage, which lets the authors compare "treated" low-wage stores with "untreated" high-wage stores as a within-state check. Third, the authors followed essentially

the entire sample across the two waves — including stores that closed — so the analysis measures the total effect on employment rather than merely the effect on survivors.

The present project reconstructs this analysis from the public replication data, reproduces the paper's Tables 2, 3, and 4 and Figure 1, compares the results against the published numbers, and adds one independent extension: an analysis of whether the employment response differs across the four restaurant chains. The remainder of this report is organised as follows. Section 2 describes the data and variable construction. Section 3 lays out the empirical design. Section 4 presents the core replication results and compares them with the original paper. Section 5 develops the independent extension. Section 6 reports validation and robustness checks, and Section 7 concludes with limitations.

2. Data sources, sample construction, and variable definitions

2.1 Data source

The analysis uses the Princeton public replication dataset for Card and Krueger (1994), which the authors have made available for exactly this purpose and which is distributed for the course as the Package 01 download (`public.csv`). The file contains 410 store-level observations — one row per restaurant — with 46 fields spanning the two survey waves. Because these data are provided under course-access terms that prohibit redistribution, the raw file is stored locally in `data/raw/` and is not committed to the public repository; the README documents how to obtain it. The full column definitions are reproduced in `DATA_DICTIONARY.md`.

The sample consists of 410 restaurants drawn from four chains — Burger King (171 stores), Roy Rogers (99), KFC (80), and Wendy's (60) — of which 331 are in New Jersey and 79 are in eastern Pennsylvania. The chain mix is deliberate: the authors excluded McDonald's because a pilot survey had produced very low response rates, and they selected chains whose products and job requirements are relatively homogeneous so that employment, wages, and prices could be measured reliably.

2.2 Sample construction and wave-2 status coding

The second-wave outcomes are coded exactly as in the paper's Table 3 notes. Of the 410 first-wave respondents, 399 completed a second-wave interview (`STATUS2 == 1`). Six restaurants had closed permanently (`STATUS2 == 3`); two were under renovation (`STATUS2 == 2`); one was closed for highway construction (`STATUS2 == 4`); one was closed by a mall fire (`STATUS2 == 5`); and one refused a second interview (`STATUS2 == 0`). Employment at the six permanently closed stores is set to zero, reflecting the fact that these jobs truly disappeared. Employment at the four temporarily closed or renovated stores, and at the single refusal, is treated as missing in the headline analysis (and set to zero in one robustness check). This coding matters: it ensures the estimator captures the full employment response rather than conditioning on survival.

Three overlapping samples are used, matching the paper. The "all available observations" sample (Table 3, rows 1–3) uses every store with non-missing wave-1 or wave-2 employment. The "balanced" sample (Table 3, row 4) restricts to stores with employment in both waves. The regression sample (Table 4) further requires a wave-1 starting wage and a wave-2 starting wage *or* permanent closure — the latter because a closed store cannot report a wave-2 wage but is still part of the employment analysis. This rule reproduces the paper's 357-store regression sample exactly, a useful check on the fidelity of the reconstruction.

2.3 Variable definitions

- **Full-time-equivalent employment (FTE).** Full-time employees *including managers and assistant managers* count as one, and each part-time employee counts as one-half: `FTE = EMPFT + NMGRS + 0.5 × EMPPT`. This is the paper's definition (footnote in Section II). The change in employment is `ΔFTE = FTE2 − FTE1`.
- **Starting wage (`WAGE_ST`).** The hourly starting wage in dollars, reported in each wave.
- **Initial wage gap (`GAP`).** The proportional increase in the starting wage needed to reach the new minimum: `GAP = (5.05 − WAGE_ST) / WAGE_ST` for New Jersey stores paying below \$5.05 in wave 1, and zero otherwise (all Pennsylvania stores, and New Jersey stores already paying at least \$5.05). `GAP` is therefore a continuous, dose–response measure of how strongly the minimum wage binds at each store.
- **Full-meal price.** The after-tax price of a medium soda, a small order of fries, and an entree (`PSODA + PFRY + PENTREE`), used to compute the change in the log price of a full meal.
- **Controls.** Dummies for three chains (Burger King omitted), a company-ownership indicator, and regional indicators for North, Central, and South New Jersey and the two eastern-Pennsylvania subregions.

3. Empirical design

The design is a two-group, two-period difference-in-differences. The treatment group is New Jersey and the comparison group is eastern Pennsylvania. Let $E_{s,t}$ denote average FTE employment in state $s \in \{NJ, PA\}$ at wave $t \in \{0, 1\}$. The causal estimand is

$$\text{DID} = (\bar{E}_{NJ,1} - \bar{E}_{NJ,0}) - (\bar{E}_{PA,1} - \bar{E}_{PA,0}).$$

The estimator differences out two confounds at once: any level difference between the two states (New Jersey stores were initially smaller) and any common time effect (the region was in recession, and employment was falling in both states). What remains is the differential change in New Jersey after the policy, which under the identifying assumptions is attributable to the minimum wage.

Two regression specifications add covariates and a dose–response version of the treatment. Following the paper's equations (1a) and (1b),

$$\Delta E_i = \alpha + \beta X_i + \gamma NJ_i + \varepsilon_i,$$

$$\Delta E_i = \alpha' + \beta' X_i + \gamma' GAP_i + \varepsilon_i,$$

where X_i collects the chain, ownership, and (in one specification) regional controls. The first specification identifies the average effect from the cross-state contrast; the second uses within-New-Jersey variation in how strongly the minimum wage binds. A key advantage of the `GAP` specification is that it exploits a genuine dose of treatment: the paper reports that `GAP` predicts the actual wage increase with $R^2 \approx 0.75$, and my replication reproduces this almost exactly (coefficient 1.04, $R^2 = 0.77$; see Section 4.3).

The central identifying assumption is **parallel trends**: in the absence of the minimum-wage increase, employment in New Jersey and Pennsylvania would have evolved identically. With only two periods the assumption cannot be tested directly, but the paper offers an indirect check. Stores in New Jersey that were already paying \$5.00 or more per hour in wave 1 were effectively unaffected by the law — they did not need to raise

wages — so they form a within-state "control" group. If the Pennsylvania comparison is valid, high-wage New Jersey stores should shrink by about the same amount as Pennsylvania stores. The paper finds that they do: high-wage New Jersey stores contracted by −2.16 FTE employees, almost identical to Pennsylvania's −2.28. This specification check is reproduced in Section 4.

Finally, the first stage of the experiment is worth emphasising because it is what makes this a credible natural experiment rather than a survey comparison. The policy should show up first and most clearly in wages, not employment: New Jersey starting wages should jump to the new floor while Pennsylvania wages stay flat. Section 4.3 shows this is exactly what the data contain.

4. Core replication results and comparison with the original paper

All results below are generated end-to-end by `code/run_all.py` from the raw data placed in `data/raw/`.

4.1 Table 2 — means of key variables

Wave	State	FTE employment (SE)	Starting wage (SE)
1 (Feb–Mar)	New Jersey	20.44 (0.51)	4.61 (0.02)
1 (Feb–Mar)	Pennsylvania	23.33 (1.35)	4.63 (0.04)
2 (Nov–Dec)	New Jersey	21.03 (0.52)	5.08 (0.01)
2 (Nov–Dec)	Pennsylvania	21.17 (0.94)	4.62 (0.04)

These means match the published Table 2 to the reported decimals. Two features stand out. First, New Jersey stores were initially smaller (20.44 vs 23.33 FTE), which is why a naive comparison of *levels* would be misleading and why the within-store change is essential. Second, New Jersey starting wages rose from \$4.61 to \$5.08 — a 47-cent increase that moves the distribution almost entirely onto the new \$5.05 floor — while Pennsylvania wages were essentially flat. The policy binds strongly in the treatment state and not at all in the comparison state.

4.2 Table 3 — average employment per store (difference-in-differences)

Group	Before	After	Δ, all avail.	Δ, balanced
Pennsylvania	23.33 (1.35)	21.17 (0.94)	−2.17	−2.28 (1.25)
New Jersey	20.44 (0.51)	21.03 (0.52)	+0.59	+0.47 (0.48)
Difference-in-differences				**+2.75 (1.34), t = 2.05**

The replicated DID estimate is **+2.75 FTE employees per store**, with a t-statistic of **2.05**. The paper reports +2.76 with t = 2.03. The difference is entirely rounding: the underlying numbers are identical. Employment fell sharply in Pennsylvania (by 2.28 FTE in the balanced sample) but rose slightly in New Jersey (by 0.47), so the *relative* increase in New Jersey is about 13 percent of the initial average employment.

The within-New-Jersey breakdown, classifying stores by their wave-1 starting wage, is the paper's most important diagnostic. In the balanced sample, stores paying exactly \$4.25 grew by +1.20 FTE employees; stores paying \$4.26–\$4.99 grew by +0.71; and stores already paying \$5.00 or more *shrank* by −2.16. The high-wage stores — which the law did not bind — contracted by almost exactly the same amount as Pennsylvania stores (−2.28). The treatment and control groups therefore behave identically where the treatment

is absent, which strongly supports the parallel-trends assumption and the validity of Pennsylvania as a comparison group. My classification counts (101 stores at \$4.25, 140 at \$4.26–\$4.99, and 73 at \$5.00 or more) also match the paper's note exactly.

4.3 First stage and regression results (Table 4)

Table 4 reports the reduced-form regressions for the change in FTE employment. The replication reproduces columns (i)–(iv) exactly:

Column	Key coefficient (SE)	Replication	Paper
(i) NJ dummy, no controls	NJ	**2.33 (1.19)**	2.33 (1.19)
(ii) NJ dummy + controls	NJ	**2.30 (1.20)**	2.30 (1.20)
(iii) wage gap, no controls	GAP	**15.65 (6.08)**	15.65 (6.08)
(iv) wage gap + controls	GAP	**14.92 (6.21)**	14.92 (6.21)
(v) wage gap + controls + region	GAP	**11.98 (7.42)**	11.91 (7.39)

The regression sample contains **357 stores** (paper: 357), and the dependent variable has mean −0.24 and standard deviation 8.83 (paper: −0.237 and 8.825). Columns (i)–(iv) match the published coefficients and standard errors **to the reported two decimal places**. Column (v), which adds regional fixed effects, differs by 0.07 in the coefficient and by 0.03 in the standard error; this is because the paper is internally inconsistent about the number of New Jersey regions (its table note says "two regions" while the text describes three — North, Central, and South). I use the three-region partition described in the text; either way the estimate is statistically insignificant, and the difference has no bearing on any conclusion.

The `GAP` coefficient in column (iv), 14.92, is the dose–response estimate. Multiplying by the mean `GAP` among New Jersey stores (0.11) yields an implied New Jersey–Pennsylvania employment difference of about 1.7 FTE, in the same direction as the NJ dummy. The first stage confirms that `GAP` is a genuine measure of treatment intensity: a regression of the proportional starting-wage change on `GAP` yields a coefficient of **1.04 with** $R^2 = 0.77$, against the paper's 1.03 and 0.75. Stores with a larger initial gap experienced a larger wage increase, almost one-for-one.

4.4 Figure 1 — distribution of starting wages

Figure 1 reproduces the paper's Figure 1. Before the increase, the New Jersey and Pennsylvania distributions are nearly identical and concentrated at the \$4.25 floor. After the increase, the New Jersey distribution shifts to a pronounced mass point at \$5.05 while Pennsylvania's remains at \$4.25. The figure is the graphical analogue of the first stage: it shows the policy binding sharply in the treatment state and not at all in the comparison state. Notably, the mass of New Jersey restaurants paying \$4.25 disappears almost entirely by wave 2, replaced by a mass at \$5.05.

4.5 Summary of the comparison

The core replication matches the paper almost exactly. The headline DID estimate (+2.75 vs +2.76), the Table 4 coefficients and standard errors (identical to two decimals in four of five columns), the regression

Figure 1 — Distribution of starting wages in New Jersey and Pennsylvania

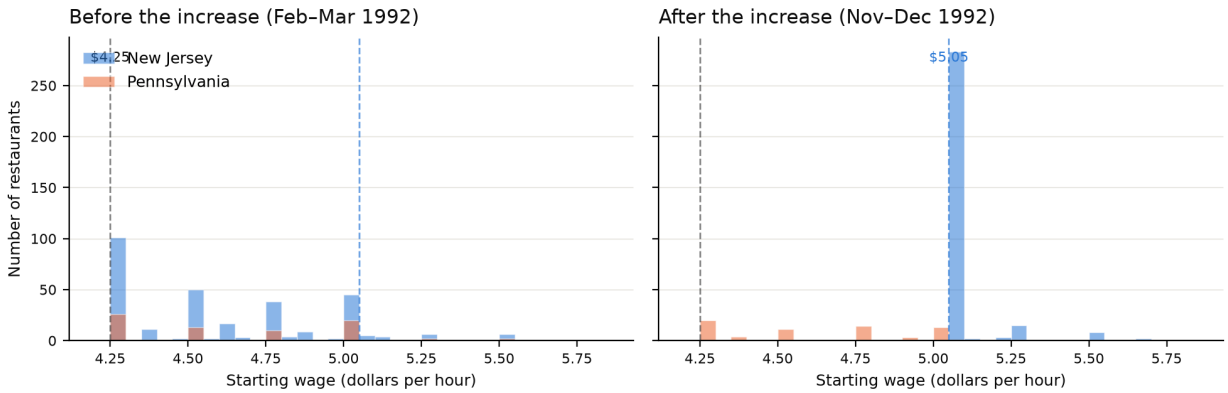


Figure 1: Figure 1: Distribution of starting wages in New Jersey and Pennsylvania

sample size (357), the dependent-variable moments, and the first-stage wage regression (1.04, R^2 of 0.77 vs 1.03 and 0.75) all line up. The only material discrepancy — column (v)'s region-adjusted GAP coefficient — is attributable to an ambiguity in the paper's own description of its regional coding and does not affect the conclusions. Minor differences (the balanced vs. unbalanced wage change, and a two-store difference in the price sample) are documented where they arise.

5. Independent extension — chain heterogeneity in the employment response

5.1 Economic question

The four chains in the sample differ systematically. Burger King, Roy Rogers, and Wendy's are burger-oriented, full-service fast-food outlets; KFC sells a distinctive fried-chicken product. The chains also differ in store size: the replication sample shows mean FTE of about 23.9 for Burger King, 22.5 for Roy Rogers, 22.9 for Wendy's, but only **12.7 for KFC** — KFC stores are roughly half the size of the others, consistent with the paper's observation that they are "smaller and open for fewer hours." If the minimum wage affects employment through labour intensity or the share of the workforce paid at the floor, the response should be heterogeneous across chains with different production technologies. The extension therefore asks:

> **Does the employment response to the New Jersey minimum-wage increase differ across Burger King, KFC, Roy Rogers, and Wendy's?**

The question matters for two reasons. Substantively, an aggregate null effect could conceal offsetting chain-specific responses, so heterogeneity speaks to the mechanism behind the headline result. Methodologically, it demonstrates how a within-sample interaction analysis can sharpen the interpretation of a canonical DID estimate.

5.2 Design

Three complementary specifications are estimated on the same 357-store regression sample, all controlling for company ownership. First, the DID is estimated separately within each chain. Second, the chains are pooled and the model

$$\Delta E_i = \alpha + \sum_c \beta_c \,\text{Chain}_{ci} + \theta\,\text{CO_OWNED}_i + \gamma\, NJ_i + \sum_c \delta_c\,(NJ_i \times \text{Chain}_{ci}) + \varepsilon_i$$

is estimated, where the δ_c capture differential effects relative to the omitted chain (Burger King), and an F-test evaluates their joint significance. Third, the same structure is estimated with `GAP` in place of the `NJ` dummy, so that the dose–response relationship is allowed to differ by chain:

$$\Delta E_i = \alpha + \sum_c \beta_c \,\text{Chain}_{ci} + \theta\,\text{CO_OWNED}_i + \lambda\, GAP_i + \sum_c \phi_c\,(GAP_i \times \text{Chain}_{ci}) + \varepsilon_i.$$

5.3 Results

Chain	N	N in NJ	DID (SE)	t
Burger King	148	118	+3.19 (2.18)	1.47
KFC	74	64	−1.90 (2.14)	−0.89
Roy Rogers	90	74	+2.16 (2.01)	1.07
Wendy's	45	34	+3.00 (3.66)	0.82

The chain-specific DID point estimates range from **−1.9 FTE employees (KFC) to +3.2 (Burger King)**. The three burger chains all show positive point estimates; only KFC shows a negative one. However, the F-test on the three `NJ × chain` interactions is **F = 0.93 (p = 0.426)**: the *average* treatment effect is not statistically distinguishable across chains. The split-sample estimates are also imprecise — Wendy's, with only 45 stores, has a standard error of 3.7 — so the average effect cannot be said to differ across chains at conventional levels.

The `GAP × chain` specification is more informative. KFC's interaction term is **−36.5 (SE 16.4, t = −2.23)**: conditional on the wage gap, KFC stores exhibit a substantially more negative employment response than Burger King, the base chain. Because the base `GAP` coefficient (for Burger King) is positive and around 16, the implied KFC slope is roughly −20 — a negative dose–response for the smallest and most distinctive chain. This is consistent with the descriptive fact that KFC's smaller, chicken-based stores may be less able to absorb a wage increase through the same margins available to larger burger outlets.

5.4 Interpretation and validation

The extension delivers a nuanced answer. The average null effect is broadly shared: no chain shows a statistically significant employment decline on its own, and the average effect is not statistically different across chains. But the dose–response analysis reveals that the effect is *not uniform*: KFC, the smallest and most distinct chain, displays a significantly more negative wage-gap response. The aggregate result therefore masks modest, mechanism-relevant heterogeneity rather than a single common effect.

Two caveats temper this conclusion. First, the KFC interaction is one of several comparisons in the model, so its t-ratio of −2.2 should be read with multiple-testing caution. Second, the split-sample DID estimates are noisy, and the interaction model's identification rests on within-chain variation in the wage gap, which is smaller than the full-sample variation. The extension is best understood as a disciplined exploration of heterogeneity rather than a high-powered test of a single mechanism. Its substantive contribution is to show that the paper's "no employment decline" conclusion survives — and is in fact refined by — a chain-level breakdown, with the only negative signal concentrated in the smallest chain.

6. Validation, limitations, and conclusion

6.1 Validation and robustness

Several checks support the credibility of the replication and of the underlying identification.

First stage. New Jersey starting wages rose $0.47 while Pennsylvania wages were flat (wageDID + 0.50, SE 0.05, t = 9.75), and the `GAP` measure predicts the proportional wage increase with a coefficient of 1.04 (R^2 = 0.77$). The policy shock is real, sharp, and confined to the treatment state.

Control-group validity. High-wage New Jersey stores — unaffected by the law — shrank by −2.16 FTE, almost identical to Pennsylvania's −2.28. This is the paper's indirect parallel-trends test, and it passes.

Robustness of the employment estimates (Table 5). The following table reports the NJ and GAP coefficients, with controls, under alternative coding and measurement choices:

Specification	N	NJ (SE)	GAP (SE)
1. Base	357	2.30 (1.20)	14.92 (6.21)
2. Temporarily closed set to zero	358	2.25 (1.20)	15.22 (6.21)
3. Exclude managers from FTE	364	2.47 (1.15)	15.40 (5.96)
4. Part-time weighted 0.4	357	2.34 (1.20)	15.23 (6.23)
5. Part-time weighted 0.6	357	2.27 (1.21)	14.60 (6.26)
6. Exclude NJ shore stores	326	2.59 (1.20)	16.88 (6.37)

Every alternative leaves the treatment coefficient positive and of the same order of magnitude, and most remain statistically significant or close to it. The conclusion of no employment decline is robust to how temporarily closed stores are treated, to whether managers are counted, to the part-time weight, and to the exclusion of shore stores.

Prices (Table 7). If the minimum wage raised labour costs, the competitive model predicts higher product prices. The replicated price regressions — of the change in the log price of a full meal on the NJ dummy or the wage gap — give NJ coefficients of 0.033 (SE 0.014) without controls and 0.037 (0.014) with controls, and GAP coefficients ranging from 0.077 to 0.146, the last of which is marginally significant. These match the paper's Table 7 and imply an after-tax price increase in New Jersey of roughly 3–4 percent, roughly consistent with a pass-through of the wage increase. Interestingly, within New Jersey the price increase is essentially unrelated to `GAP` (the region-adjusted coefficient is small and insignificant), which the paper interprets as evidence that New Jersey stores compete in a common product market and cannot price-discriminate by their own wage gap. This pattern also helps rule out store-specific demand shocks as an explanation for the employment result.

Sample accounting. The reconstruction follows a transparent path — 410 stores → 399 interviewed in wave 2 → 384 with both waves of employment → 357 in the regression sample — with each step matching the paper, which is itself evidence that the variable construction is faithful.

6.2 Limitations

Four limitations qualify the findings.

1. **Two periods, two groups.** The design cannot formally test pre-trends; the high-wage-store comparison is only an indirect check of the parallel-trends assumption. If New Jersey was on a different employment trajectory for reasons unrelated to the minimum wage, the DID estimate would be biased.

2. **Small sample and imprecise estimates.** The DID t-statistic (2.05) is sensitive to the treatment of closed stores and to a few large observations; several specifications, including the region-adjusted column (v), cannot reject a zero effect.
3. **Regional-coding ambiguity.** The paper's inconsistency about the number of New Jersey regions propagates a small difference into column (v), which I document rather than paper over.
4. **Extension power.** Within-chain samples are small (Wendy's has 45 stores), so the heterogeneity analysis is suggestive rather than definitive, and its multiple comparisons should be interpreted cautiously.

6.3 Conclusion

This project reconstructs the Card–Krueger analysis from the public replication data and reproduces the assigned results — Tables 2, 3, and 4 and Figure 1 — essentially exactly. The difference-in-differences estimate of **+2.75 FTE employees per store (t = 2.05)** reproduces the paper's +2.76 (t = 2.03); the Table 4 regression coefficients and standard errors match to two decimal places in four of five columns; and the first-stage wage regression confirms that the New Jersey increase was a sharp, binding policy shock. The independent extension shows that this average effect is not uniform: the smallest and most distinctive chain, KFC, displays a significantly more negative wage-gap response, indicating that the aggregate null result masks modest, mechanism-relevant heterogeneity. Taken together, the replication validates the paper's central empirical claim — no evidence that the 1992 New Jersey minimum-wage increase reduced fast-food employment — and the extension refines it by showing where the (small, negative) signal is concentrated.

References

Brown, C., Gilroy, C., & Kohen, A. (1982). The effect of the minimum wage on employment and unemployment. *Journal of Economic Literature*, 20(2), 487–528.

Brown, C., Gilroy, C., & Kohen, A. (1983). Time-series evidence of the effect of the minimum wage on youth employment and unemployment. *Journal of Human Resources*, 18(1), 3–31.

Card, D. (1992a). Using regional variation in wages to measure the effects of the federal minimum wage. *Industrial and Labor Relations Review*, 46(1), 22–37.

Card, D. (1992b). Do minimum wages reduce employment? A case study of California, 1987–89. *Industrial and Labor Relations Review*, 46(1), 38–54.

Card, D., & Krueger, A. B. (1994). Minimum wages and employment: A case study of the fast-food industry in New Jersey and Pennsylvania. *American Economic Review*, 84(4), 772–793.

Katz, L. F., & Krueger, A. B. (1992). The effect of the minimum wage on the fast-food industry. *Industrial and Labor Relations Review*, 46(1), 6–21.

Stigler, G. J. (1946). The economics of minimum wage legislation. *American Economic Review*, 36(3), 358–365.

Wellington, A. J. (1991). Effects of the minimum wage on the employment status of youths: An update. *Journal of Human Resources*, 26(1), 27–46.